Mistral-Nemo RTBT vs. DPO: theory vs. experiment

199 profiles selected for DPO training

Candidate probabilities

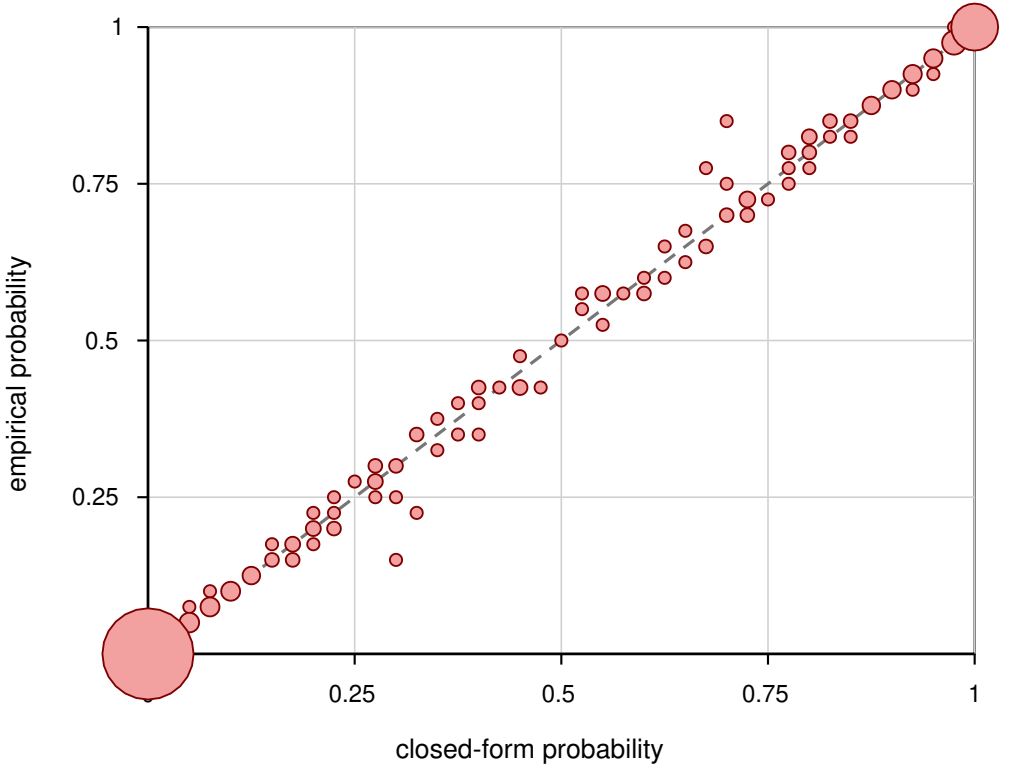

Largest bins: n=499 at (0, 0), n=108 at (1, 1); r=1.000; MAE=0.0027; SD|err|=0.0105

Profile-level discrepancy

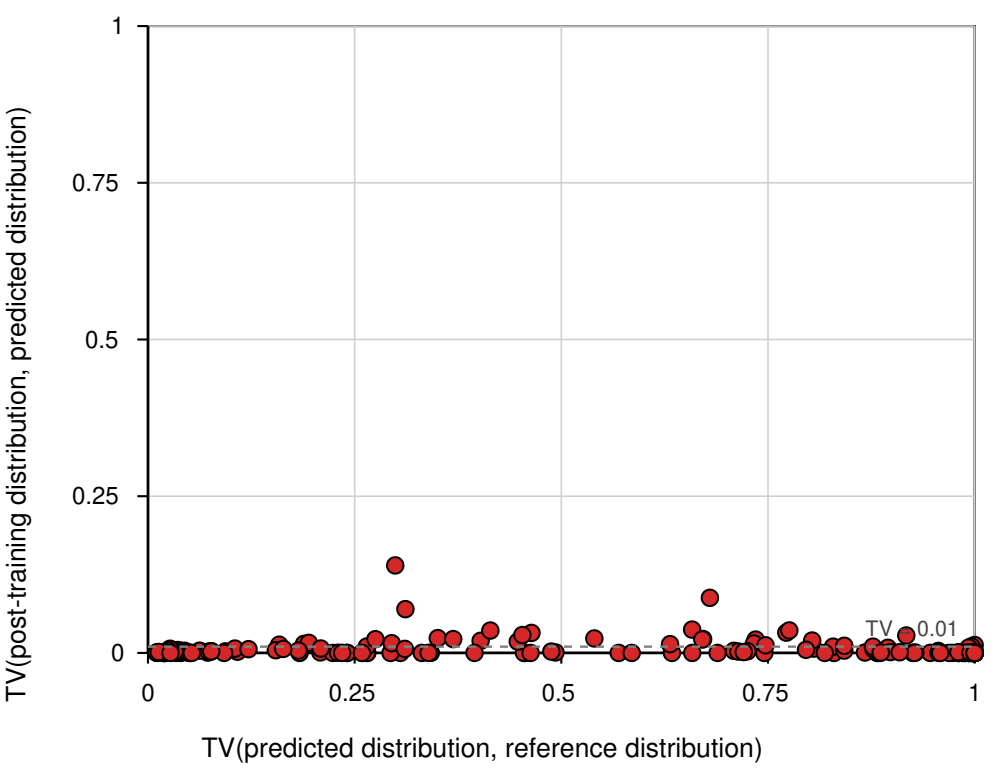

median TV=0.00023; mean TV=0.00547; SD TV=0.01440; top match=199/199