Qwen3 RTBT vs. DPO: theory vs. experiment

167 profiles selected for DPO training

Candidate probabilities

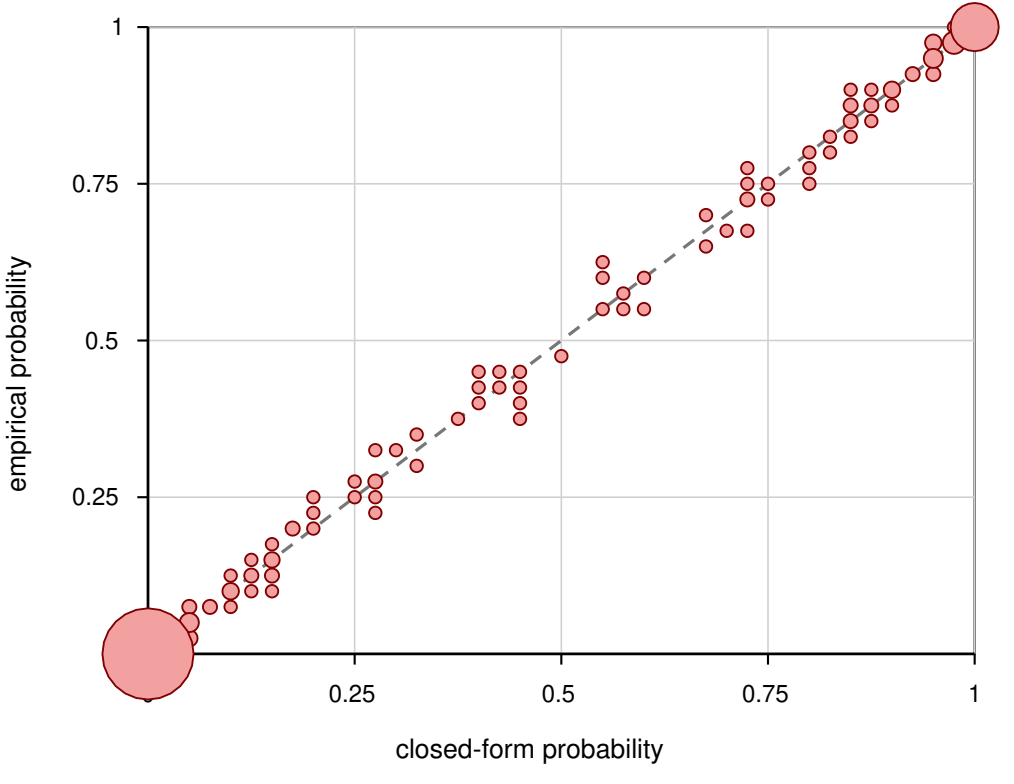

Largest bins: n=431 at (0, 0), n=99 at (1, 1); r=1.000; MAE=0.0026; SD|err|=0.0082

Profile-level discrepancy

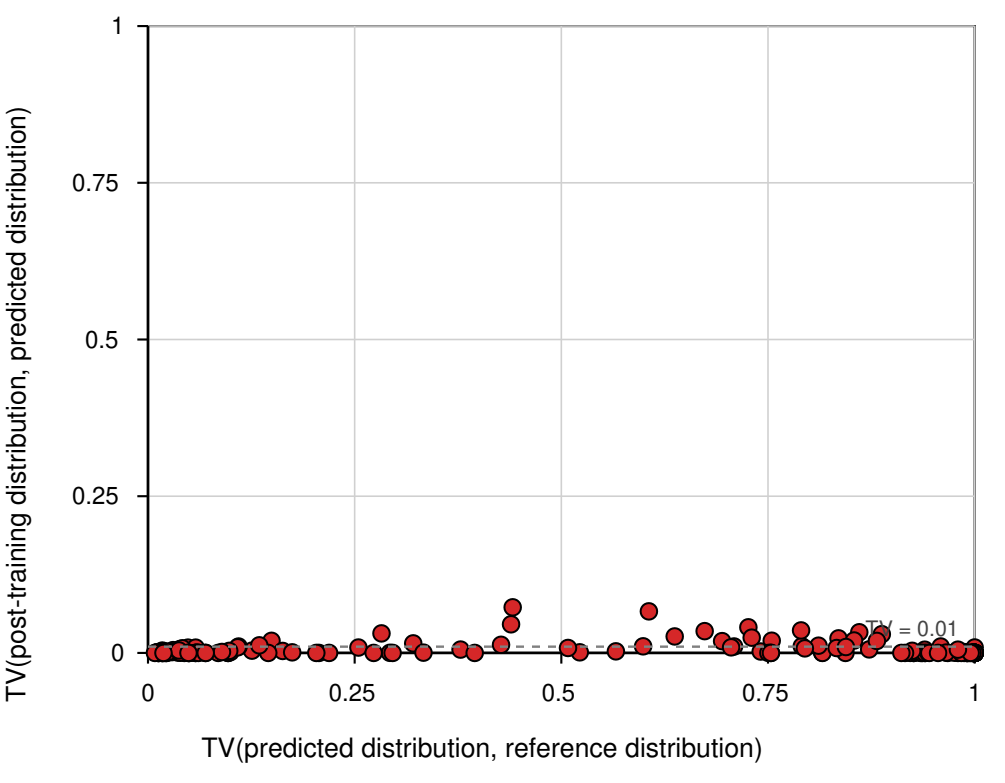

median TV=0.00011; mean TV=0.00510; SD TV=0.01110; top match=167/167