Qwen2.5 RTBT vs. DPO: theory vs. experiment

154 profiles selected for DPO training

Candidate probabilities

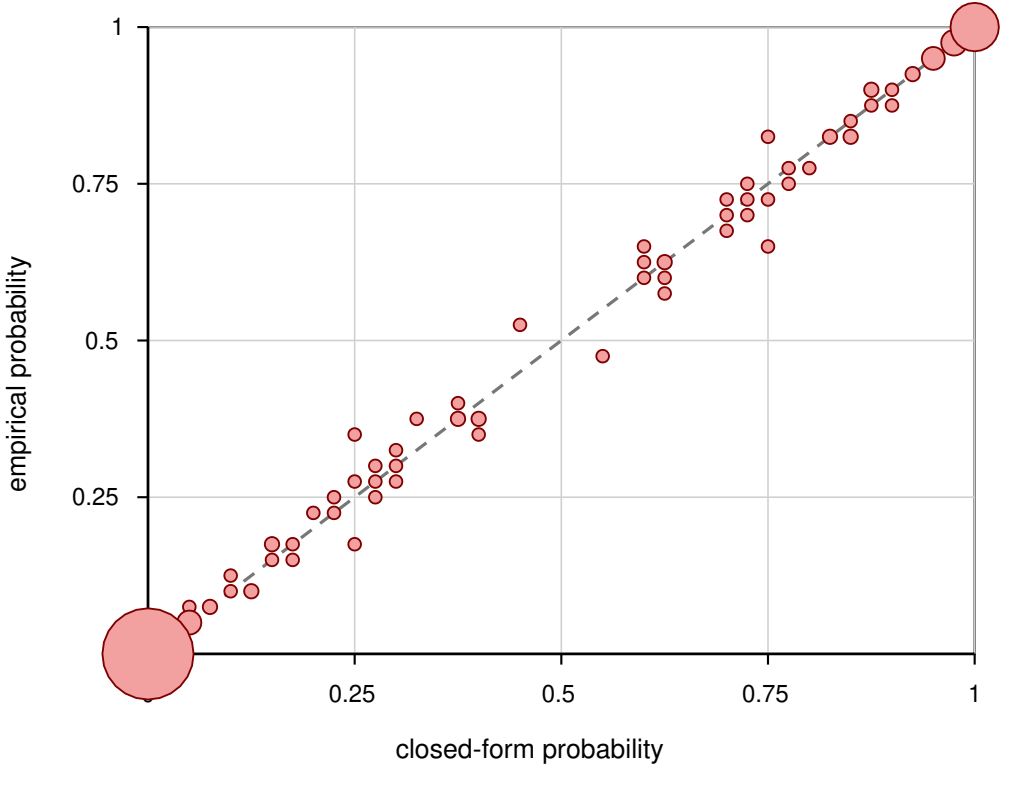

Largest bins: n=396 at (0, 0), n=92 at (1, 1); r=1.000; MAE=0.0023; SD|err|=0.0094

Profile-level discrepancy

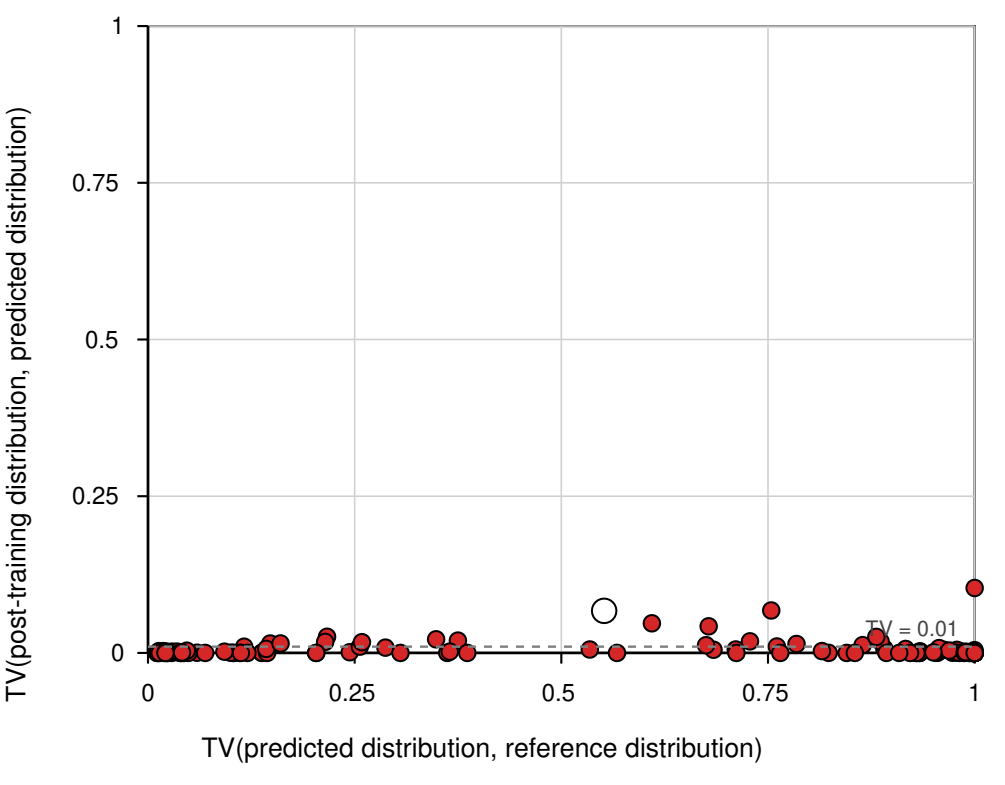

median TV=0.00009; mean TV=0.00465; SD TV=0.01294; top match=153/154