Qwen2.5 RTBT vs. joint DPO: theory vs. experiment

154 profiles; one shared LoRA adapter trained jointly

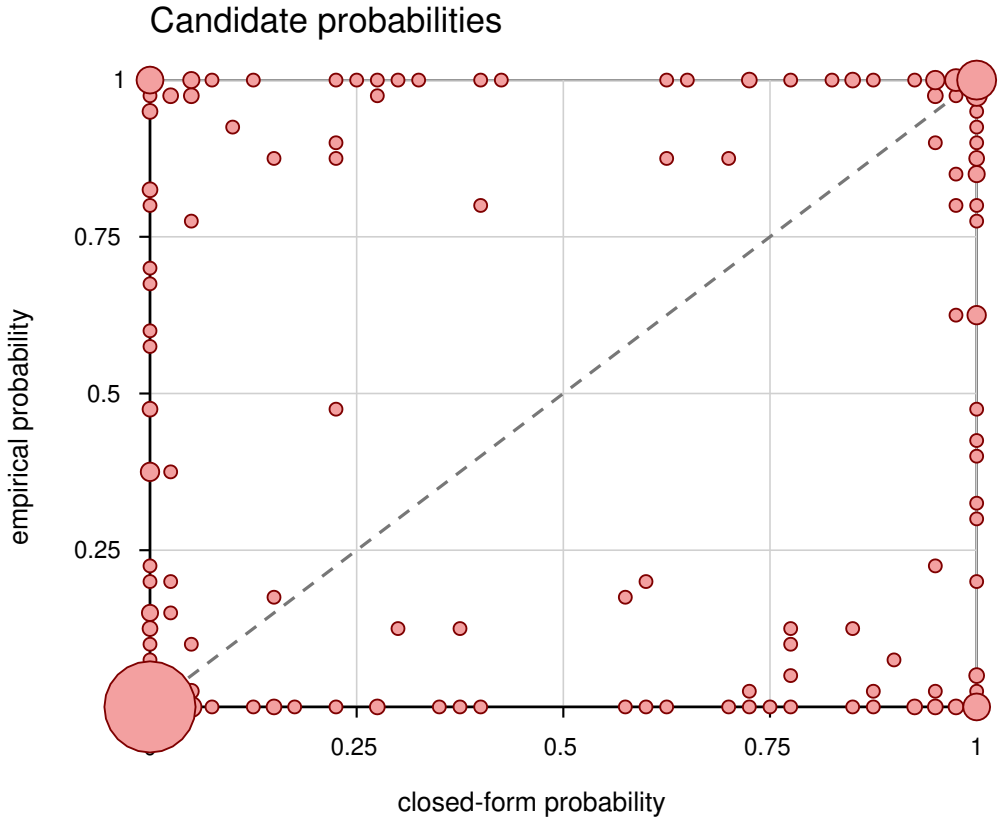

Largest bins: n=346 at (0, 0), n=47 at (1, 1); r=0.613; MAE=0.1665; SD|err|=0.3180

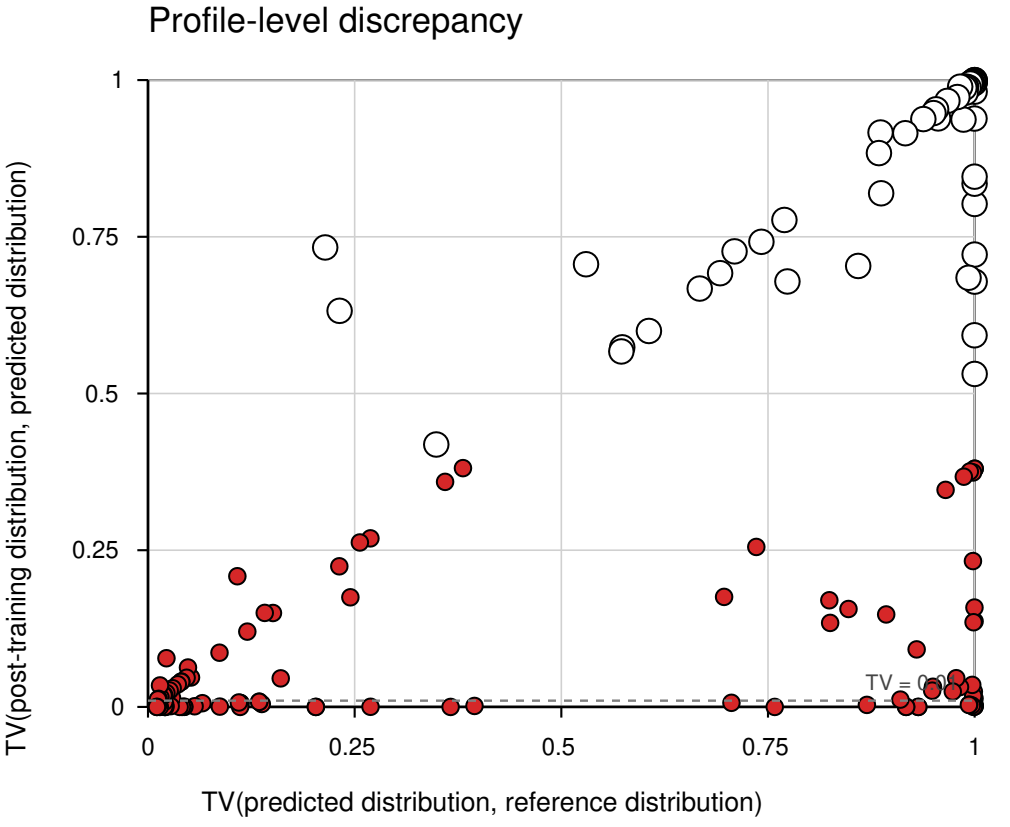

median TV=0.12709; mean TV=0.33291; SD TV=0.38739; top match=102/154